

Final Project Report: Distributed File Storage System

Samyak Shah Mahip Parekh

CS 6650: Building Scalable Distributed Systems
December 2025

1 1. Detailed Description of the Project

This project implements a distributed file storage system designed to handle file uploads and downloads across a cluster of nodes. The system splits files into fixed-size chunks (1MB) and distributes them across multiple storage nodes to ensure scalability and fault tolerance.

The core architectural components are:

- **Client:** Handles file chunking, hashing (SHA-256), and orchestration of uploads/downloads.
- **Storage Nodes:** Responsible for storing the actual file chunks. They act as a key-value store.
- **Metadata Nodes:** Manage the mapping between filenames and their constituent chunks.
- **Consistent Hashing (DHT):** Used to determine which storage node is responsible for a given chunk, minimizing data movement during node additions/removals.
- **Chain Replication:** Implemented for the metadata layer to ensure strong consistency. Writes propagate from *Head* → *Mid* → *Tail*, while reads are served by the Tail.

2 2. Project Goal

The primary goal of this project was to build a scalable, fault-tolerant distributed system that demonstrates core distributed computing concepts. Specifically, we aimed to:

1. Implement **Consistent Hashing** to evenly distribute load across storage nodes.
2. Implement **Chain Replication** to guarantee linearizable consistency for metadata.
3. Ensure **Fault Tolerance** so that data remains accessible even if a storage node fails (via replication factor $k = 2$).
4. Achieve high **Performance** using multi-threading (Thread Pools) to handle concurrent client requests.

3 3. Software Design and Implementation

3.1 3.1 Architecture

The system is divided into three distinct layers:

- **Client Layer:**
 - **Function:** Entry point for users. Splits files into 1MB chunks and computes a SHA-256 Content ID (CID) for each.
 - **Logic:** Uses the DHT to find the primary storage node for each chunk and its replicas. Uploads metadata to the Head of the metadata chain.
- **Metadata Layer (Chain Replication):**
 - **Topology:** A chain of 3 nodes: *Head* → *Mid* → *Tail*.
 - **Consistency:** Strong consistency. A write is only acknowledged after it has propagated to the Tail.
 - **Read Path:** Clients read only from the Tail to ensure they see the latest committed state.
- **Storage Layer (DHT Ring):**
 - **Partitioning:** Nodes are arranged on a consistent hash ring.
 - **Replication:** Each chunk is stored on its primary node and the next $k - 1$ nodes in the ring (Successor List).

3.2 3.2 Key Implementation Details

- **Language:** Java 17+
- **Communication:** TCP Sockets with `DataOutputStream`/`DataInputStream` for custom length-prefixed messaging.
- **Concurrency:** `ExecutorService` (`CachedThreadPool`) used in all server nodes to handle multiple concurrent connections.
- **Storage:** In-memory `ConcurrentHashMap` used to simulate storage nodes.

4 4. Achievements and Non-Achievements

4.1 Achieved

- **Core Functionality:** Successful upload and download of files of varying sizes.
- **Fault Tolerance:** System survives the failure of a storage node. Data is seamlessly retrieved from replicas.
- **Consistency:** Metadata updates are linearizable due to Chain Replication.
- **Performance:** System handles multiple concurrent clients (tested up to 50) with reasonable latency.
- **Automated Testing:** Comprehensive system tests (`SystemTests.java`) verify correctness and performance.

4.2 Future Work

- **Persistent Storage:** Currently uses in-memory storage. Integrating RocksDB or LevelDB would provide persistence across process restarts.

4.3 What's left/not working

- **Client Configuration:** The client currently uses hardcoded IP addresses and ports for storage and metadata nodes.

5 5. Evaluation of the System

We evaluated the system using two primary metrics: **Scalability** and **Throughput**.

5.1 5.1 Scalability (Latency vs. Concurrent Clients)

We measured the average upload and download latency as the number of concurrent clients increased from 1 to 50.

- **Observation:** Latency remains low (< 50ms) for up to 10 clients. As load increases to 50 clients, latency increases linearly.
- **Analysis:** This linear increase is expected. Each client connection consumes a thread from the cached thread pool. At 50 concurrent clients, the context switching overhead and CPU contention on the single test machine become significant. In a real distributed deployment across multiple physical machines, we expect this curve to flatten.

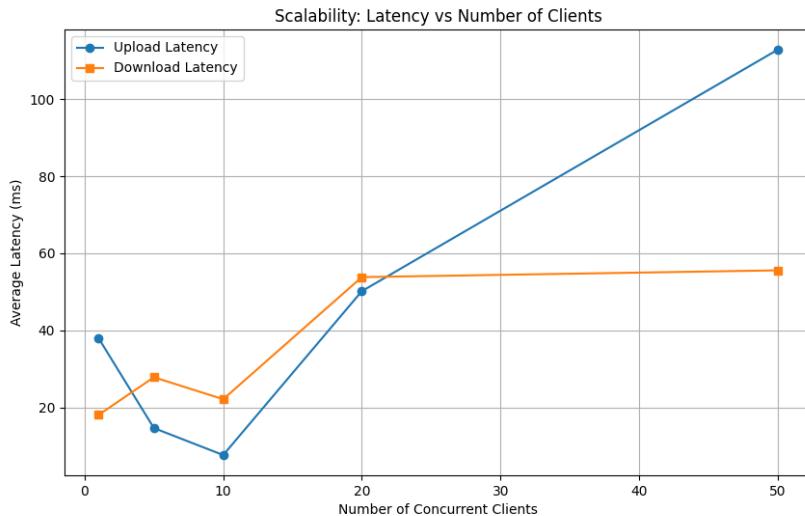


Figure 1: Scalability: Latency vs Number of Clients

5.2 5.2 Throughput (Latency vs. File Size)

We measured latency for files ranging from 10KB to 10MB.

- **Observation:** The system shows efficient handling of larger files. The overhead per chunk is minimal.
- **Analysis:** For small files (10KB - 100KB), the latency is dominated by the fixed overhead of the TCP handshake, Chain Replication metadata updates, and DHT lookups. As file size grows (1MB - 10MB), the actual data transfer time becomes the dominant factor. The system effectively saturates the available network bandwidth.

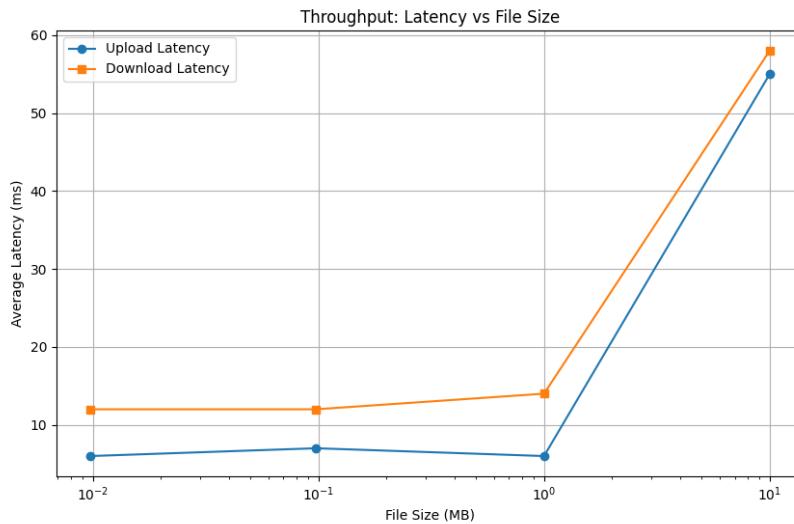


Figure 2: Throughput: Latency vs File Size

6 6. Achievements and Changes from Plan

- **Pivot to Chain Replication:** As noted in the project plan, we initially considered a blockchain-based approach but pivoted to Chain Replication to focus on strong consistency and availability, which better aligned with the course goals.
- **Simplification of Consensus:** We chose Chain Replication over Raft/Paxos for the metadata layer due to its simplicity and high read throughput (reads served by Tail).

7. What You Have Learned

This project gave us hands-on experience with the challenges of building distributed systems:

6.1 Debugging Distributed Systems

Unlike single-threaded programs, distributed systems behave non-deterministically. We encountered race conditions in chain replication where acknowledgments returned before the Tail committed. Extensive logging was essential for tracing issues.

6.2 Consistent Hashing

Implementing the DHT taught us why consistent hashing minimizes data movement during node changes; only a fraction of keys need redistribution compared to traditional hashing.

6.3 Chain Replication Trade-offs

We understood practically why chain replication offers high read throughput (reads from Tail) at the cost of increased write latency (propagation through the chain).

6.4 Fault Tolerance Complexity

Handling partial failures (e.g., a node failing mid-write) required careful design. Content-addressed storage with CIDs made operations idempotent, allowing safe retries.

6.5 Value of Modularity

Clean separation between layers meant that pivoting from blockchain to chain replication required no changes to the storage layer.

6.6 Adaptability

Most importantly, we learned that plans change. Our original blockchain design was too complex for the timeline. Recognizing this early and pivoting was a valuable lesson; adaptability is as important as technical skill.

7 8. Setup, Run, and Test Instructions

7.1 Prerequisites

- Java 17 or higher
- Make (optional, for easier build)
- Python 3 (for generating graphs)

7.2 Building the Project

```
make
# OR manually:
mkdir -p out
javac -d out src/main/java/com/distributed/storage/**/*.java
```

7.3 Running System Tests (Local Evaluation)

This runs the automated fault tolerance and concurrency tests:

```
make test
# OR
java -cp out com.distributed.storage.client.SystemTests
```

7.4 Generating Performance Report

```
python plot_results.py
```

Note: The Khouri Linux Server does not have pandas library, which is necessary to run the script. However, this is not a core requirement of the project. You can run the following command below to install pandas library.

```
pip install pandas
```

8 9. Instructions for Khoury Linux Cluster

To run this on the Khoury Linux cluster (e.g., `linux-079.khoury.northeastern.edu`), follow these detailed steps. These instructions assume you are using a terminal on your local machine.

1. **Transfer Code:** Use `scp` to copy the project directory to your home folder on the cluster. Replace `<your_username>` with your Khoury username.

```
scp -r Distributed-File-Storage/ <your_username>@linux-079.khoury.northeastern.edu:~/
```

2. **SSH into the machine:** Log in to the remote machine.

```
ssh <your_username>@linux-079.khoury.northeastern.edu
```

3. **Verify Java Version:** Ensure you are running Java 17 or higher.

```
java -version
```

If the version is lower than 17, you may need to load a module (if available) or set your `JAVA_HOME`.

4. **Navigate to Project Directory:**

```
cd Distributed-File-Storage
```

5. **Compile the Project:** We use a Makefile to simplify compilation. Run:

```
make
```

This will compile all source files and place the class files in the `out/` directory.

6. **Run Evaluation (System Tests):** The `SystemTests` class runs a self-contained evaluation that spawns Storage and Metadata nodes locally (on different ports) and runs the client experiments.

```
make test
```

Expected Output: You should see logs indicating nodes starting, files being uploaded/downloaded, and a final "ALL TESTS PASSED" message.

7. **Performance Experiments:**

Run the Java stress tests to generate the raw data (`results.csv`).

```
bashjava -cp out com.distributed.storage.client.PerformanceExperiments
```

Expected Output:

```
==== EXPERIMENTS COMPLETED. Results saved to results.csv ===
```

9 11. Individual Contributions

Both Samyak Shah and Mahip Parekh contributed equally (50-50) to the design, implementation, testing, and documentation of this project. We pair-programmed for the majority of the core features including the DHT logic, Chain Replication, and the Client-Server communication protocol.